

Project Report

Team no: 8

Gagan BV | Bhoomika L | Shruthi P

Title: Comic Crafter AI

Problem Statement:

Creation of An AI Based Comic Generator Capable of Running Locally on Edge devices that generates a comic Style Story based on the Input.

Project Overview:

The **AI Comic Generator** is a creative tool that uses artificial intelligence to automatically generate a 4-panel comic strip based on a user-provided story prompt. It combines text generation (using Mistral-7B or GPT-2) and image generation (using Stable Diffusion) to create a cohesive comic with a retro cartoon aesthetic.

This project showcases how AI can assist in creative storytelling and visual art, blending natural language processing (NLP) and generative AI into an engaging comic-making experience.

Key Features:

- Story Generation – The AI writes a short story based on the user's input.

- Panel Splitting – The story is divided into four parts, each serving as a comic panel's caption.
- Retro Cartoon Art – Stable Diffusion generates 1940s-style cartoon images matching each panel's text.
- Automatic Layout – The system assembles the comic strip with captions and images.
- User-Friendly Interface – Built with Gradio, allowing easy interaction in a web browser.

Technical Highlights:

- **Hugging Face Transformers** (for text generation)
- **Stable Diffusion** (for image generation)
- **4-bit Quantization** (for efficient GPU usage)
- **Gradio UI** (for interactive web deployment)

Steps to Run on Google Colab:

- Open Google Colab: Go to Google Colab.
- Create a New Notebook: Click on "File" > "New Notebook".
- Install Required Libraries: Copy and paste the following code into a cell and run it:

```
>>!pip install -q -r requirements.txt
```

- Copy the Code: Copy the entire code provided in this repository into a new cell in the Colab notebook.
- Set Hugging Face Token: Replace the hf_token variable in the code with your Hugging Face token. You can obtain a token by creating an account on Hugging Face.

- Run the Code: Execute the cell containing the code. After a few moments, the Gradio interface will launch.
- Interact with the App: Enter a story prompt and click "Generate Comic" to see the generated comic strip.

Steps to Run Locally:

- Install Python: Ensure you have Python 3.7 or higher installed. You can download it from python.org.
- Set Up a Virtual Environment (Optional): It is recommended to create a virtual environment to manage dependencies.

```
>>python -m venv comic_crafter_env
```

```
>>source comic_crafter_env/bin/activate # On Windows use
`comic_crafter_env\Scripts\activate`
```

- Install Required Libraries: Use pip to install the required libraries:

```
>>pip install -r requirements.txt
```

- Copy the Code: Copy the entire code provided in this repository into a Python file (e.g., comic_crafter.py).
- Set Hugging Face Token: Replace the hf_token variable in the code with your Hugging Face token.
- Run the Code: Execute the Python file (terminal / Command prompt):

```
python comic_crafter.py
```

- Access the App: After running the code, a local server will start. Open the provided URL (usually <http://127.0.0.1:7860>) in your web browser to access the Gradio interface.

- Interact with the App: Enter a story prompt and click "Generate Comic" to see the generated comic strip.

Code:

```
# Install required libraries

#!pip install -q transformers diffusers torch matplotlib bitsandbytes gradio

# Import libraries

from transformers import AutoModelForCausalLM, AutoTokenizer,
BitsAndBytesConfig

from diffusers import StableDiffusionPipeline

import torch

import gc

from PIL import Image, ImageDraw, ImageFont

import gradio as gr

import random

# Verify GPU availability

device = "cuda" if torch.cuda.is_available() else "cpu"

print(f"Using device: {device}")

# Hugging Face token for gated models (replace with your token)

hf_token = "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"

# Configure 4-bit quantization for text generation (only if CUDA is available)

quantization_config = None

if device == "cuda":

    try:

        quantization_config = BitsAndBytesConfig(

            load_in_4bit=True,
```

```

        bnb_4bit_use_double_quant=True,
        bnb_4bit_quant_type="nf4",
        bnb_4bit_compute_dtype=torch.float16
    )
except Exception as e:
    print(f"Quantization not supported: {e}. Falling back to CPU.")
    device = "cpu"
    quantization_config = None

# Load the text generation model (Mistral 7B)
try:
    text_model_name = "mistralai/Mistral-7B-v0.1"
    text_tokenizer = AutoTokenizer.from_pretrained(text_model_name,
        use_auth_token=hf_token)
    text_model = AutoModelForCausalLM.from_pretrained(
        text_model_name,
        quantization_config=quantization_config,
        device_map="auto",
        use_auth_token=hf_token
    )
except Exception as e:
    print(f"Failed to load text model: {e}")
    # Fallback to a smaller model
    text_model_name = "gpt2"
    text_tokenizer = AutoTokenizer.from_pretrained(text_model_name)
    text_model = AutoModelForCausalLM.from_pretrained(text_model_name).to(device)

```

```
# Load the image generation model (Stable Diffusion)
```

```
try:
```

```
    image_model_name = "runwayml/stable-diffusion-v1-5"
```

```
    image_pipe
```

```
=
```

```
StableDiffusionPipeline.from_pretrained(image_model_name,  
torch_dtype=torch.float16)
```

```
    image_pipe = image_pipe.to(device)
```

```
except Exception as e:
```

```
    print(f"Failed to load image model: {e}")
```

```
    raise RuntimeError("Image model could not be loaded. Please check your  
setup.")
```

```
# Function to generate text (story or dialogue)
```

```
def generate_text(prompt):
```

```
    try:
```

```
        with torch.no_grad():
```

```
            inputs = text_tokenizer(prompt, return_tensors="pt").to(device)
```

```
            outputs = text_model.generate(  
                **inputs,
```

```
                max_length=300,
```

```
                num_return_sequences=1,
```

```
                temperature=0.7,
```

```
                top_k=50,
```

```
                top_p=0.9,
```

```
                do_sample=True,
```

```
                pad_token_id=text_tokenizer.eos_token_id
```

```
            )
```

```

        generated_text = text_tokenizer.decode(outputs[0],
skip_special_tokens=True)

    # Clean up

    del inputs, outputs

    gc.collect()

    torch.cuda.empty_cache()

    return generated_text

except Exception as e:

    print(f"Error generating text: {e}")

    return f"Error generating story: {str(e)}"

```

Function to generate a retro cartoon-style image

```
def generate_retro_cartoon_image(prompt):
```

```
    try:
```

```
        retro_cartoon_prompt = (
```

```
            f"{prompt}, classic 1940s cartoon style, hand-drawn animation, "
```

```
            "bold outlines, vibrant colors, cel-shaded, exaggerated expressions, "
```

```
            "vintage animation, retro cartoon"
```

```
        )
```

```
    with torch.no_grad():
```

```
        image = image_pipe(
```

```
            retro_cartoon_prompt,
```

```
            height=512,
```

```
            width=512,
```

```
            num_inference_steps=30 # More steps for better quality
```

```

        ).images[0]

    return image

except Exception as e:

    print(f"Error generating image: {e}")

    # Create error placeholder image

    img = Image.new('RGB', (512, 512), color=(255, 200, 200))

    draw = ImageDraw.Draw(img)

    draw.text((50, 250), "Failed to generate image", fill=(0, 0, 0))

    return img

# Function to split story into 4 parts for panels
def split_story(story):

    sentences = [s.strip() for s in story.split('.') if s.strip()]

    if len(sentences) < 4:

        # If not enough sentences, duplicate some

        sentences = sentences * (4 // len(sentences) + 1)

    # Distribute sentences across 4 panels

    panel_texts = []

    for i in range(4):

        start = i * len(sentences) // 4

        end = (i + 1) * len(sentences) // 4

        panel_text = ' '.join(sentences[start:end]) + ' '

        panel_texts.append(panel_text)

    return panel_texts

# Function to create comic panel with text
def create_comic_panel(text, image, panel_size=(512, 512)):

```


try:

```
panel = Image.new("RGB", panel_size, "white")
draw = ImageDraw.Draw(panel)
# Resize and paste image (top 75% of panel)
img_height = int(panel_size[1] * 0.75)
panel.paste(image.resize((panel_size[0], img_height)), (0, 0))
# Add text (bottom 25%)
font = ImageFont.load_default()
text_position = (10, img_height + 10)
draw.text(text_position, text, fill="black", font=font)
return panel
```

except Exception as e:

```
print(f"Error creating panel: {e}")
error_img = Image.new('RGB', panel_size, color=(255, 200, 200))
draw = ImageDraw.Draw(error_img)
draw.text((50, 250), "Panel creation error", fill=(0, 0, 0))
return error_img
```

Main generation function

def generate_comic(story_prompt):

try:

```
# Generate story
story = generate_text(f"Write a short story about: {story_prompt}")
# Split into 4 parts
panel_texts = split_story(story)
# Generate panels
```

```

panels = []
for i, text in enumerate(panel_texts):
    image = generate_retro_cartoon_image(text)
    panel = create_comic_panel(text, image)
    panels.append(panel)
return story, *panels
except Exception as e:
    print(f"Error in comic generation: {e}")
    error_img = Image.new('RGB', (512, 512), color=(255, 200, 200))
    draw = ImageDraw.Draw(error_img)
    draw.text((50, 250), "Generation error", fill=(0, 0, 0))
    return f"Error: {str(e)}", error_img, error_img, error_img, error_img

```

Custom CSS for styling

```

css = """
.gradio-container {
    max-width: 1200px !important;
}
.panel-container {
    display: flex;
    flex-wrap: wrap;
    justify-content: center;
    gap: 10px;
    margin-top: 20px;
}

```

```
.panel{  
    border: 2px solid #ddd;  
    border-radius: 8px;  
    box-shadow: 0 4px 6px rgba(0,0,0,0.1);  
}
```

```
.story-box {  
    background: #f8f9fa;  
    padding: 15px;  
    border-radius: 8px;  
    margin-bottom: 20px;  
}
```

```
""""
```

Gradio interface

with gr.Blocks(css=css, title="AI Comic Generator") as demo:

```
    gr.Markdown("""
```

```
    # 🎨 AI Comic Generator
```

```
    Enter a story idea below and the AI will generate a 4-panel comic strip!
```

```
""")
```

```
    with gr.Row():
```

```
        with gr.Column():
```

```
            story_prompt = gr.Textbox(  
                label="Story Prompt",  
                placeholder="e.g., 'A robot who wants to be a chef',  
                lines=3  
            )
```

```

generate_btn = gr.Button("Generate Comic", variant="primary")

with gr.Row():
    story_output = gr.Textbox(
        label="Generated Story",
        interactive=False,
        elem_classes=["story-box"]
    )
with gr.Row(elem_classes=["panel-container"]):
    panel_outputs = []
    for i in range(4):
        panel_outputs.append(
            gr.Image(
                label=f"Panel {i+1}",
                elem_classes=["panel"],
                width=512,
                height=512
            )
        )
generate_btn.click(
    fn=generate_comic,
    inputs=story_prompt,
    outputs=[story_output] + panel_outputs
)
demo.launch(share=True) # Launch the app

```

Code Description:

Model Loading & Setup

- Text Generation Model (Mistral-7B/GPT-2):
 - Loads a large language model (LLM) for generating stories.
 - Uses 4-bit quantization (if GPU is available) to optimize memory usage.
 - Falls back to GPT-2 if Mistral-7B fails.
- Image Generation Model (Stable Diffusion):
 - Loads Stable Diffusion v1.5 for generating retro cartoon-style images.
 - Uses FP16 precision for faster GPU inference.

Text Generation

- Takes a user-provided prompt (e.g., "A robot who wants to be a chef").
- Generates a short story (300 tokens max) with controlled randomness (temperature, top-k, top-p sampling).
- Cleans up GPU memory afterward to avoid crashes.

Image Generation

- Converts each story segment into a retro cartoon-style image by modifying the prompt (e.g., adding "1940s cartoon style").
- Ensures 512x512 resolution with 30 inference steps for better quality.
- Returns a placeholder image if generation fails.

Comic Panel Creation

- Splits the story into 4 segments (for 4 comic panels).
- If the story is too short, repeats sentences to fill panels.
- Combines text and images into panels:
 - Top 75% = AI-generated image.
 - Bottom 25% = Story text (using PIL for drawing).

Gradio Interface

- Provides a user-friendly web UI with:
 - A text input box for the story idea.
 - A "Generate Comic" button to trigger the process.
 - A text output box showing the full story.
 - Four image panels displaying the comic.
- Uses custom CSS for a polished look.

Error Handling & Fallbacks

- Gracefully handles failures (e.g., if models don't load).
- Falls back to simpler models (GPT-2 if Mistral fails).
- Displays error images/text instead of crashing.

Performance Optimization

- Clears GPU cache after each generation to prevent memory leaks.
- Uses quantization (4-bit) to reduce VRAM usage.
- Batches operations efficiently (text → split → images → panels).

Final Output

- Returns a 4-panel comic strip with:
 - A coherent AI-generated story.
 - Vintage cartoon-style illustrations.
 - Text captions below each image.